

Performance Profiling in Python

Network vs CPU Bottlenecks Analysis

Karen Cardiel Olea
Universidad Politécnica de Yucatán
Email: 22090239@upy.edu.mx

Abstract—This technical report presents a case study on Python profiling techniques to identify and distinguish between I/O-bound and CPU-bound performance bottlenecks. Two scenarios were analyzed, downloading weather station data from NOAA servers (I/O-bound) and computing distances between geographic stations (CPU-bound). Using tools such as cProfile, SnakeViz, and line_profiler, to demonstrate how proper profiling enables targeted optimization strategies.

I Section 2.1: Network / I-O Profiling

I-A What I Did

I executed a Python script that downloads weather data from NOAA servers for multiple stations and years. I compared two versions:

- **Without cache** (`load.py`): downloads data every time from NOAA servers.
- **With cache** (`load_cache.py`): checks if file exists before downloading using `os.path.exists()`, reusing previously downloaded files.

Commands executed:

First run (without cache):

```
python3 -m cProfile -s cumulative load.py
01044099999,02293099999 2021 2021 > profile.txt
```

Second run (with cache):

```
python3 -m cProfile -s cumulative load_cache.py
01044099999,02293099999 2021 2021 > profile_cache.txt
```

I-B Results

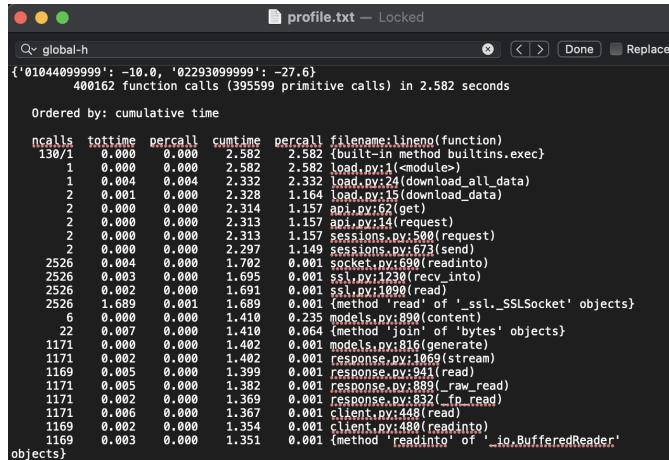


Fig. 1. cProfile results without caching - showing high `requests.get` time

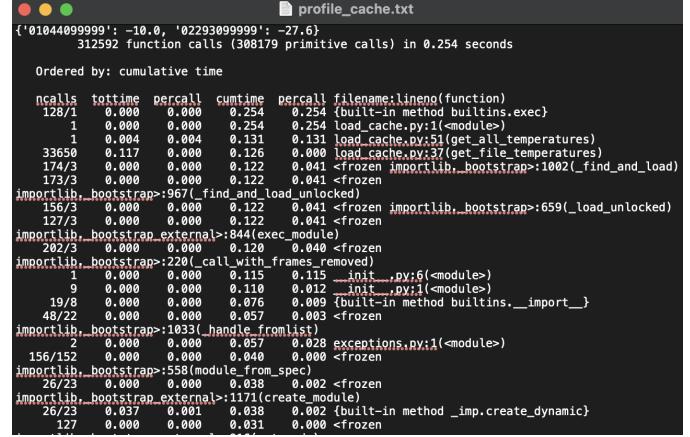


Fig. 2. cProfile results with caching - showing reduced network calls

I-C Observations and Interpretation

Most of the time was spent in:

- `requests.get()` - HTTP requests to NOAA servers
- SSL socket reads and network I/O operations

After implementing caching:

- Execution time decreased significantly on subsequent runs
- Network calls reduced to zero for already-downloaded files

This demonstrates a clear **I/O-bound process**, where performance depends primarily on network speed and latency rather than computational capacity. The simple file existence check (`os.path.exists()`) eliminated redundant downloads and dramatically improved performance.

II Section 2.2: CPU Profiling

II-A What I Did

I cloned the code repository from the book:

```
git clone https://github.com/tiagoantao/python-performance
```

The relevant code is located in `02-python/sec2-cpu/`.

I computed pairwise distances between weather stations using geographic coordinates with the Haversine formula. This operation has $O(n^2)$ complexity, making it computationally intensive.

Tools used:

- cProfile: High-level performance overview
- SnakeViz: Visualization of profiling results
- line_profiler: Per-line performance analysis

Execution steps:

Step 1: Generate profiling file

```
python3 -m cProfile -o distance_cache.prof
distance_cache.py
```

Step 2: Visualize with SnakeViz

```
python3 -m snakeviz distance_cache.prof
```

Step 3: Line profiling

```
python3 -m kernprof -l lprofile_distance_cache.py
python3 -m line_profiler lprofile_distance_cache.py.
lprof
```

II-B Results

```
Command executed 1 hr ago and took 9 mins
karen@Karens-MacBook-Air sec2-cpu % python3 -m cProfile -o distance_cache.prof distance_cache.py
```

Fig. 3. Terminal after running cProfile

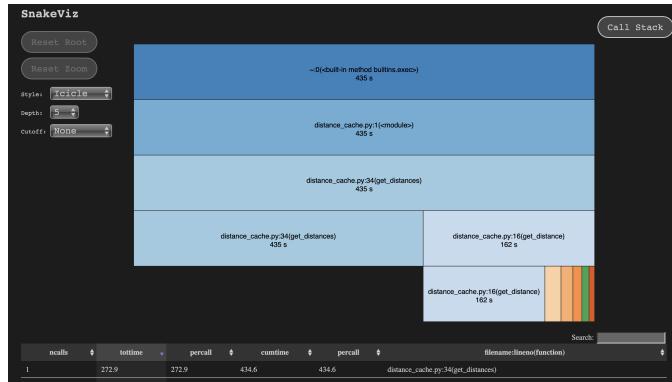


Fig. 4. SnakeViz visualization showing time distribution

```
karen@Karens-MacBook-Air sec2-cpu % python3 -m line_profiler lprofile_distance_cache.py.lprof
Timer unit: 1e-06 s
Total time: 142.888 s
File: lprofile_distance_cache.py
Function: get_distance at line 16
Line #  Hits   Time  Per Hit  % Time  Line Contents
=====
16          @profile
17          def get_distance(p1, p2):
18              lat1, lon1 = p1
19              lat2, lon2 = p2
20
21      36888866    9547766.0    0.3    6.7
22      36888866   825697.0    0.2    5.8
23      36888865   5848311.0    0.2    4.1
24  73761731  22410676.0    0.3   15.7
25  118642596  30198785.0    0.3   21.1
26  73761730  17839899.0    0.2   12.5
27
28  36888865  14471921.0    0.4   10.1
29  36888865  5990663.0    0.2    4.2
30  36888865  7043624.0    0.2    4.9
31
32  36888865  7059291.0    0.2    4.9
                                return dist
```

Fig. 5. Line profiler output showing bottleneck lines

Key metrics:

- Total execution time: 142.888 seconds
- Function analyzed: `get_distance()`
- Total calls per line: ~36.8 million
- Algorithm complexity: $O(n^2)$

TABLE I
MOST EXPENSIVE LINES IN GET_DISTANCE()

Line	Operation	% Time
24	<code>math.cos(radians(lat1)) * math.cos(...)</code>	21.1%
23	<code>math.sin(lat_dist / 2) * math.sin(...)</code>	15.7%
25	<code>math.sin(lon_dist / 2) * math.sin(...)</code>	12.5%
28	<code>math.atan2(...)</code>	10.1%

II-C Observations and Interpretation

Majority of execution time spent in mathematical operations:

- Trigonometric functions (`math.sin`, `math.cos`)
- Repeated `math.radians()` conversions
- `math.atan2()` calculations

Lines 23-25 combined account for nearly 50% of total execution time. The bottleneck is heavy computation repeated millions of times. Each operation is fast (approximately 0.2-0.4 microseconds), but 36+ million calls create significant overhead.

This confirms a **CPU-bound problem**. Unlike Section 2.1, this slowdown is not related to waiting for external resources, but to:

- Pure computation with intensive mathematical operations
- Algorithmic complexity: $O(n^2)$ means calculations grow quadratically
- Pattern: *Small cost × Massive repetition = Major bottleneck*

III Conclusions

Key findings:

- Profiling tools are essential to identify bottleneck origins before optimization
- I/O-bound tasks benefit from caching and reducing redundant operations (Section 2.1)
- CPU-bound tasks require algorithmic or implementation-level optimization (Section 2.2)
- Visual tools (SnakeViz) and line profiling provide targeted insights
- Scaling up is mostly about swapping slow $O(n^2)$ logic for faster ($O(n)$) solutions.”

References

- [1] T. Antao, *Fast Python: High Performance Techniques for Large Datasets*. Manning Publications, 2023, ch. 2, pp. 18-28.